
MEGState: Phoneme Decoding from Magnetoencephalography Signals

Anonymous Author(s)

Affiliation

Address

email

Abstract

Decoding linguistically meaningful representations from non-invasive neural recordings remains a central challenge in neural speech decoding. Among available neuroimaging modalities, magnetoencephalography (MEG) provides a safe and repeatable means of mapping speech-related cortical dynamics, yet its low signal-to-noise ratio and high temporal dimensionality continue to hinder robust decoding. In this work, we introduce MEGState, a novel architecture for phoneme decoding from MEG signals that captures fine-grained cortical responses evoked by auditory stimuli. Extensive experiments on the LibriBrain dataset demonstrate that MEGState consistently surpasses baseline model across multiple evaluation metrics. These findings highlight the potential of MEG-based phoneme decoding as a scalable pathway toward non-invasive brain–computer interfaces for speech.

1 Introduction

Decoding speech representations from brain activity holds considerable promise for restoring communication in individuals with paralysis or severe speech impairments [Moses et al., 2021]. Recent advances in invasive brain–computer interfaces [Nagashima et al., 2025, Suzuki et al., 2025] have enabled continuous speech reconstruction from intracranial recordings, achieving word error rates below 5% for vocabularies exceeding 100,000 words [Card et al., 2024]. However, their reliance on neurosurgical implantation limits scalability and clinical feasibility. In contrast, non-invasive approaches such as magnetoencephalography (MEG) offer a safe and repeatable alternative for probing speech-related neural activity. Nevertheless, decoding linguistically meaningful information from MEG remains challenging due to its low signal-to-noise ratio, high temporal resolution, and sparse neural representations [Yang et al., 2024b,a].

In this work, we introduce MEGState, a novel architecture designed for phoneme classification from MEG signals. The model integrates two complementary components: (i) a Multi-Resolution Convolution module that captures fine-grained temporal dynamics of phoneme-evoked cortical responses, and (ii) a Sensor-wise SSM that captures long-range temporal dependencies across individual sensors. This design effectively mitigates the challenges of MEG’s sparsity and high sampling rate, allowing the model to capture both local and global neural dynamics. Comprehensive experiments on the LibriBrain dataset demonstrate that MEGState consistently surpasses baseline method across multiple evaluation metrics.

31 2 Method

32 2.1 Preliminaries

33 **State space models.** Recent progress in state space models (SSMs) [Gu et al., 2022, Dao and Gu, 34 2024] shows that they can outperform prevailing architectures, most notably Transformers, on a wide 35 range of sequence modeling tasks. Grounded in control theory [Kalman, 1960], SSMs provide a

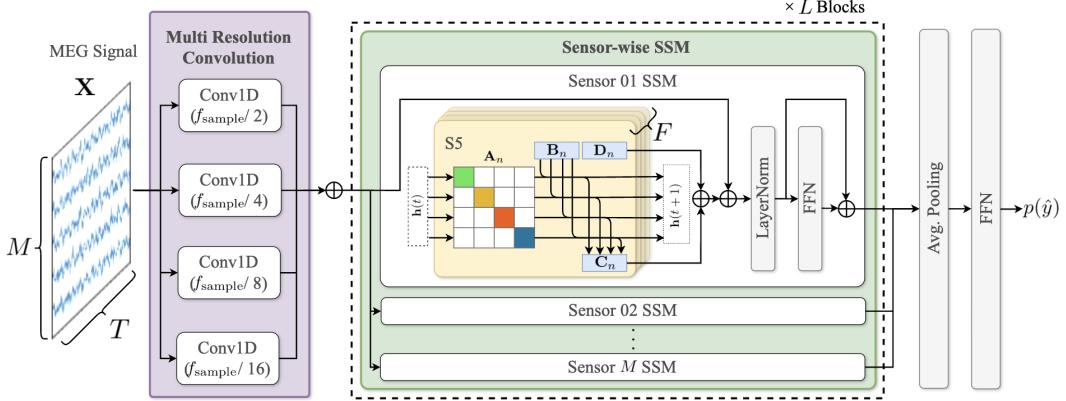


Figure 1: Model architecture of the proposed MEGState. Given a MEG signal, Mutli Resolution Convolution module extracts local temporal dependencies while Sensor-wise SSM models global spatial and temporal dependencies, respectively.

36 mapping from inputs $\mathbf{x}(t) \in \mathbb{R}^P$ to outputs $\mathbf{y}(t) \in \mathbb{R}^P$ through latent states $\mathbf{h}(t) \in \mathbb{R}^Q$ governed by

$$\frac{d\mathbf{h}(t)}{dt} = \mathbf{A}\mathbf{h}(t) + \mathbf{B}\mathbf{x}(t), \quad \mathbf{y}(t) = \mathbf{C}\mathbf{h}(t) + \mathbf{D}\mathbf{x}(t), \quad (1)$$

37 where $\mathbf{A} \in \mathbb{R}^{Q \times Q}$ is the state matrix and $\mathbf{B} \in \mathbb{R}^{Q \times P}$, $\mathbf{C} \in \mathbb{R}^{P \times Q}$, $\mathbf{D} \in \mathbb{R}^{P \times P}$ are input/output
38 projections. Among SSM variants, S5 [Smith et al., 2023] has proved especially effective for
39 modeling continuous signals. S5 sets \mathbf{A} to the HiPPO-N matrix Gu et al. [2022] to capture long-
40 range temporal dependencies. Since HiPPO-N is real symmetric, it admits the diagonalization
41 $\mathbf{A} = \mathbf{V}\Lambda\mathbf{V}^{-1}$, which yields the decoupled form

$$\frac{d\tilde{\mathbf{h}}(t)}{dt} = \Lambda\tilde{\mathbf{h}}(t) + \tilde{\mathbf{B}}\mathbf{x}(t), \quad \mathbf{y}(t) = \tilde{\mathbf{C}}\tilde{\mathbf{h}}(t) + \tilde{\mathbf{D}}\mathbf{x}(t), \quad (2)$$

42 where $\tilde{\mathbf{h}}(t) = \mathbf{V}^{-1}\mathbf{h}(t)$, $\tilde{\mathbf{B}} = \mathbf{V}^{-1}\mathbf{B}$, and $\tilde{\mathbf{C}} = \mathbf{C}\mathbf{V}$. Moreover, introducing a timescale vector
43 $\Delta \in \mathbb{R}^{+Q}$ and applying zero-order hold discretization Zhang and Chong [2007] gives the recurrence

$$\tilde{\mathbf{h}}_t = \bar{\Lambda}\tilde{\mathbf{h}}_{t-1} + \bar{\mathbf{B}}\mathbf{x}_t, \quad \mathbf{y}_t = \bar{\mathbf{C}}\tilde{\mathbf{h}}_t + \bar{\mathbf{D}}\mathbf{x}_t, \quad (3)$$

44 where $\bar{\Lambda} = \exp(\Lambda\Delta)$, $\bar{\mathbf{B}} = \Lambda^{-1}(\bar{\Lambda} - \mathbf{I})\tilde{\mathbf{B}}$, $\bar{\mathbf{C}} = \tilde{\mathbf{C}}$, $\bar{\mathbf{D}} = \mathbf{D}$. In practice, we take \mathbf{D} to be diagonal
45 and learn $\text{diag}(\Lambda)$, $\tilde{\mathbf{B}}$, $\tilde{\mathbf{C}}$, $\text{diag}(\mathbf{D})$, and Δ .

46 2.2 Model Architecture

47 The overall architecture of the proposed model is depicted in Figure 1. It comprises two primary
48 modules: Multi-Resolution Convolution and Sensor-wise SSM. Given a MEG sample $\mathbf{X} \in \mathbb{R}^{M \times T}$,
49 where M and T represent the number of sensors and the sequence length, respectively, the network
50 proceeds as follows.

51 First, the Multi-Resolution Convolution module extract fine-grained local temporal structures re-
52 flecting distinct cortical responses evoked by different phonemes. It consists of four parallel one-
53 dimensional convolutional layers with kernel sizes of $f_{\text{sample}}/2$, $f_{\text{sample}}/4$, $f_{\text{sample}}/8$, and $f_{\text{sample}}/16$,
54 where f_{sample} denotes the native sampling rate. The outputs from these convolutional layers are
55 concatenated to yield $\mathbf{H} \in \mathbb{R}^{F \times M \times T}$, where F represents the frequency feature dimension.

56 Next, the Sensor-wise SSM module models global temporal dependencies in a sensor-specific manner.
57 To handle the sparsity and high temporal resolution of MEG signals, we extend S5 [Smith et al.,
58 2023], a variant of SSM well suited for modeling high-dimensional multivariate time series. The
59 module comprises L hierarchically organized blocks, and the output $\tilde{\mathbf{H}}^{(l)} \in \mathbb{R}^{F \times M \times T}$ from the l -th
60 block ($l = 1, \dots, L$) is obtained as follows, where $\tilde{\mathbf{H}}^{(0)} = \mathbf{H}$:

$$\mathbf{H}'^{(l-1)} = \text{LayerNorm} \left(\text{SSM} \left(\tilde{\mathbf{H}}^{(l-1)} \right) + \tilde{\mathbf{H}}^{(l-1)} \right), \quad (4)$$

$$\tilde{\mathbf{H}}^{(l)} = \text{LayerNorm} \left(\text{FFN} \left(\mathbf{H}'^{(l-1)} \right) + \mathbf{H}'^{(l-1)} \right). \quad (5)$$

Algorithm 1 Sampling training data

Input: Dataset $(\mathcal{X}, \mathcal{Y}) = \{(\mathbf{X}_i, y_i)\}_{i=1}^N$

Output: Training sample $(\tilde{\mathbf{X}}, \tilde{y})$

Randomly sample two labels $y_1 \sim \text{Uniform}(\mathcal{Y}), y_2 \sim \text{Uniform}(\mathcal{Y})$

$\mathcal{I}_1 \leftarrow \{i \in \{1, \dots, N\} \mid y_i = y_1\}$ ▷ All indices with label y_1

$\mathcal{I}_2 \leftarrow \{i \in \{1, \dots, N\} \mid y_i = y_2\}$ ▷ All indices with label y_2

Randomly sample subsets $\mathcal{I}'_1 \subset \mathcal{I}_1, \mathcal{I}'_2 \subset \mathcal{I}_2$ s.t. $|\mathcal{I}'_1| = |\mathcal{I}'_2| = N'$

$\bar{\mathbf{X}}_1 \leftarrow \text{AVERAGE}_{i \in \mathcal{I}'_1}(\mathbf{X}_i)$

$\bar{\mathbf{X}}_2 \leftarrow \text{AVERAGE}_{i \in \mathcal{I}'_2}(\mathbf{X}_i)$

$\tilde{\mathbf{X}} \leftarrow \alpha \bar{\mathbf{X}}_1 + (1 - \alpha) \bar{\mathbf{X}}_2$ ▷ Mixup augmentation

$\tilde{y} \leftarrow \alpha y_1 + (1 - \alpha) y_2$

61 Here, $\text{SSM}(\cdot)$, $\text{LayerNorm}(\cdot)$, and $\text{FFN}(\cdot)$ denote the S5 layer, layer normalization, and feed-forward
 62 network, respectively. Subsequently, the output $\tilde{\mathbf{H}}^{(L)}$ from the Sensor-wise SSM is aggregated to
 63 produce the predicted phoneme probability $p(\hat{y})$ corresponding to \mathbf{X} as follows:

$$p(\hat{y}) = \text{FFN} \left(\text{AvgPool} \left(\tilde{\mathbf{H}}^{(L)} \right) \right), \quad (6)$$

64 where $\text{AvgPool}(\cdot)$ denotes average pooling over the temporal dimension, following the standard SSM
 65 architectures [Gu et al., 2022]. The model is trained using the cross-entropy loss function.

66 3 Experiments

67 3.1 Dataset and Pre-procesing

68 In this experiment, we used the publicly available LibriBrain dataset Özdogan et al. [2025], which
 69 contains MEG recordings of a single participant listening to audiobook narrations of Sherlock
 70 Holmes. The recordings were acquired with a MEGIN Triux™ Neo system using 306 sensors at
 71 1 kHz, yielding 52.32 hours of data. Preprocessing involved Maxwell filtering [Taulu and Simola,
 72 2006] and Signal Source Separation to remove sensor noise and external magnetic interference. A
 73 notch filter (50 Hz), Butterworth band-pass filter (0.1–125 Hz), and downsampling to 250 Hz were
 74 further applied. The dataset provides 39 ARPAbet-based phoneme labels temporally aligned with the
 75 auditory stimuli, where each label corresponds to a 0.5 s MEG segment. It was divided into training,
 76 validation, and test sets comprising 1,488,392, 11,289, and 12,051 samples, respectively. Models
 77 were trained, tuned, and evaluated on these splits, and further tested on the LibriBrain leaderboard set
 78 provided in the 2025 PNPL Competition [Landau et al., 2025].

79 3.2 Implementation Details

80 During training, we leverage smoothing and data augmentation in the sampling of the training samples
 81 to improve the signal-to-noise ratio (SNR) of MEG signals and mitigate phoneme-label imbalance.
 82 Concretely, at each sampling step within a training epoch, we independently draw two phoneme labels
 83 y_1 and y_2 uniformly from the MEG dataset $(\mathcal{X}, \mathcal{Y}) = \{(\mathbf{X}_i, y_i)\}_{i=1}^N$. For each label y_k ($k \in \{1, 2\}$),
 84 we construct the index set of all samples with that label, $\mathcal{I}_k = \{i \in \{1, \dots, N\} \mid y_i = y_k\}$,
 85 and uniformly sample a subset $\mathcal{I}'_k \subset \mathcal{I}_k$ of size N' . We then form a class-conditional prototype by
 86 averaging the corresponding MEG inputs, $\bar{\mathbf{X}}_k = \frac{1}{N'} \sum_{i \in \mathcal{I}'_k} \mathbf{X}_i$, which acts as a denoised representative
 87 and thus enhances the SNR. Furthermore, by introducing a mixing coefficient $\alpha \in [0, 1]$, we address
 88 label imbalance (see Sec. 4) via a mixup-style convex combination [Zhang et al., 2018] of both the
 89 prototypes and their labels:

$$\tilde{\mathbf{X}} = \alpha \bar{\mathbf{X}}_1 + (1 - \alpha) \bar{\mathbf{X}}_2, \quad \tilde{y} = \alpha \bar{y}_1 + (1 - \alpha) \bar{y}_2. \quad (7)$$

90 The overall sampling procedure is summarized in Algorithm 1.

91 We trained the model with the AdamW optimizer ($\beta_1 = 0.9, \beta_2 = 0.999$) and a learning rate of
 92 1.0×10^{-4} . The batch size was 32, and training proceeded for 50 epochs. For the Sensor-wise SSM
 93 module, we set the block size to $L = 2$. For sampling training data, we set the prototype size of
 94 $N' = 100$ and a mixing coefficient of $\alpha = 0.5$.

Table 1: Performance comparison on the LibriBrain [Özdogan et al., 2025] test set and leaderboard set. **Bold** values denote the best performances, while \dagger indicates statistical significance compared to the baseline method ($p < 0.05$). Multi-Resol Conv. indicates Mutli-Rresolution Convolution module.

Model	Test Set			Leaderboard Set
	Acc. [%] \uparrow	Kappa [%] \uparrow	Macro-F1 [%] \uparrow	Macro-F1 [%] \uparrow
Baseline [Özdogan et al., 2025]	38.80 \pm 2.40	45.71 \pm 0.74	34.82 \pm 1.92	—
Ours (w/o Mult-Resol. Conv.)	40.25 \pm 3.15	37.90 \pm 2.83 \dagger	34.77 \pm 1.83	—
Ours (w/o Sensor-wise SSM)	40.37 \pm 3.20	49.60 \pm 6.25	37.18 \pm 4.44	—
Ours (MEGState)	45.53\pm1.88\dagger	54.19\pm2.42\dagger	41.11\pm2.20\dagger	55.74 (68.41)

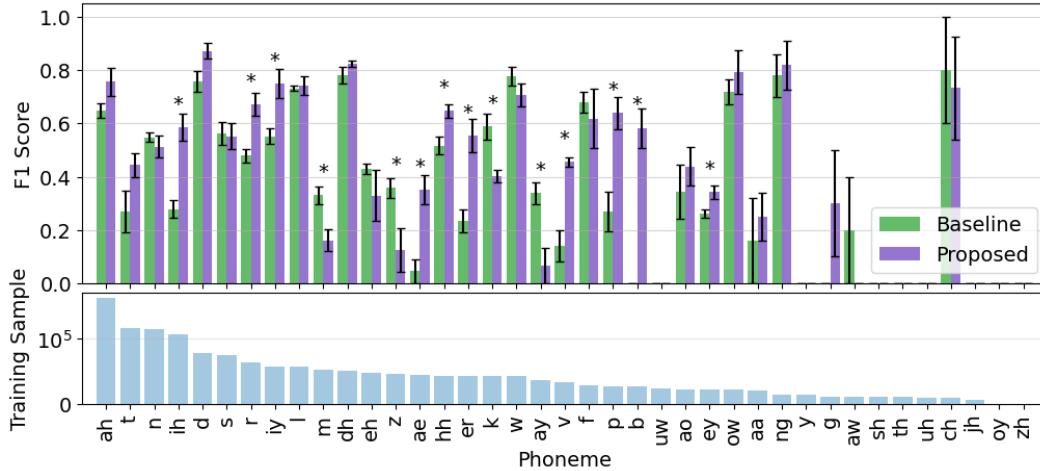


Figure 2: Quantitative comparison across phonemes. The upper panel shows the macro-F1 for each phoneme, while the lower panel indicates the number of training samples per phoneme. Error bars represent the standard error of the mean, and * denotes statistical significance ($p < 0.05$).

95 4 Results

96 Table 1 presents the quantitative comparison between the proposed and baseline methods. Values in
97 the table represent the mean and standard deviation obtained from over five distinct random seeds.

98 On the test set, the proposed method achieved balanced accuracy, Cohen’s Kappa, and macro-F1 of
99 45.53%, 54.19%, and 41.11%, respectively, surpassing the baseline method by 6.73, 8.48, and 6.29
100 points. Here, all improvements were statistically significant ($p < 0.05$). Moreover, ablation studies
101 revealed that removing either the Multi-Resolution Convolution module or the Sensor-wise SSM
102 consistently degrades performance across all metrics on the test set, indicating that both modules
103 contribute substantially to phoneme classification from MEG signals.

104 On the leaderboard set, the proposed method achieved a macro-F1 of 55.74%. Notably, the final
105 leaderboard submission employed an ensemble strategy that selected the most probable label from
106 predictions of five independently trained models, resulting in a higher macro-F1 of 68.41%.

107 Figure 2 further presents the phoneme-wise quantitative comparison. Each bar shows the macro-F1
108 score for individual phonemes, ordered by the number of training samples. As shown, the proposed
109 method outperformed the baseline in terms of macro-F1 on 19 phonemes and achieved statistically
110 significant improvements on 10 phonemes ($p < 0.05$).

111 5 Conclusion

112 In this work, we introduced MEGState, a novel architecture for phoneme decoding from MEG
113 signals. By integrating a Multi-Resolution Convolution module to capture fine-grained local temporal
114 dynamics and a Sensor-wise SSM to model long-range dependencies across individual sensors, our
115 approach effectively mitigates the challenges of MEG signal sparsity and high temporal resolution.
116 Comprehensive experiments on the LibriBrain dataset demonstrated that MEGState consistently
117 outperformed the baseline across multiple evaluation metrics.

118 **References**

- 119 Nicholas S Card, Maitreyee Wairagkar, Carrina Iacobacci, Xianda Hou, Tyler Singer-Clark, Francis R
120 Willett, Erin M Kunz, Chaofei Fan, Maryam Vahdati Nia, et al. An accurate and rapidly calibrating
121 speech neuroprosthesis. *New England Journal of Medicine*, 391(7):609–618, 2024.
- 122 Tri Dao and Albert Gu. Transformers are SSMs: Generalized Models and Efficient Algorithms
123 Through Structured State Space Duality. In *ICML*, volume 235, pages 10041–10071, 2024.
- 124 Albert Gu, Karan Goel, and Christopher Ré. Efficiently Modeling Long Sequences with Structured
125 State Spaces. In *ICLR*, 2022.
- 126 R. Kalman. A New Approach to Linear Filtering and Prediction Problems. *Journal of Basic
127 Engineering*, 82(1):35–45, 1960.
- 128 Gilad Landau, Miran Özdogan, Gereon Elvers, Francesco Mantegna, Pratik Somaiya, Dulhan Jayalath,
129 Luisa Kurth, Teyun Kwon, et al. The 2025 PNPL Competition: Speech Detection and Phoneme
130 Classification in the LibriBrain Dataset. In *NeurIPS Competition Track*, 2025.
- 131 David A Moses, Sean L Metzger, Jessie R Liu, Gopala K Anumanchipalli, Joseph G Makin, Pengfei F
132 Sun, Josh Chartier, Maximilian E Dougherty, et al. Neuroprosthesis for decoding speech in a
133 paralyzed person with anarthria. *New England Journal of Medicine*, 385(3):217–227, 2021.
- 134 Shunya Nagashima, Kanta Kaneda, Yuiga Wada, Tsumugi Iida, Misa Taguchi, Masayuki Hirata, and
135 Komei Sugiura. ECoG Dual Context Network for Brain Machine Interfaces. *IEEE Access*, 13:
136 99252–99265, 2025.
- 137 Miran Özdogan, Gilad Landau, Gereon Elvers, Dulhan Jayalath, Pratik Somaiya, Francesco Mantegna,
138 Mark Woolrich, and Oiwi Parker Jones. LibriBrain: Over 50 Hours of Within-Subject MEG to
139 Improve Speech Decoding Methods at Scale. In *NeurIPS Datasets and Benchmarks Track*, 2025.
- 140 Jimmy Smith, Andrew Warrington, and Scott Linderman. Simplified State Space Layers for Sequence
141 Modeling. In *ICLR*, 2023.
- 142 Shuntaro Suzuki, Shunya Nagashima, Masayuki Hirata, and Komei Sugiura. Cortical-SSM: A Deep
143 State Space Model for EEG and ECoG Motor Imagery Decoding. *arXiv preprint arXiv:2510.15371*,
144 2025.
- 145 Samu Taulu and Juha Simola. Spatiotemporal signal space separation method for rejecting nearby
146 interference in MEG measurements. *Physics in Medicine & Biology*, 51(7):1759, 2006.
- 147 Yiqian Yang, Yiqun Duan, Qiang Zhang, Hyejeong Jo, Jinni Zhou, Won Hee Lee, Renjing Xu, and
148 Hui Xiong. Neuspeech: Decode neural signal as speech. *arXiv preprint arXiv:2403.01748*, 2024a.
- 149 Yiqian Yang, Hyejeong Jo, Yiqun Duan, Qiang Zhang, Jinni Zhou, Won Hee Lee, Renjing Xu, and
150 Hui Xiong. Mad: Multi-alignment meg-to-text decoding. *arXiv preprint arXiv:2406.01512*, 2024b.
- 151 Hongyi Zhang, Moustapha Cisse, Yann N Dauphin, and David Lopez-Paz. mixup: Beyond Empirical
152 Risk Minimization. In *ICLR*, 2018.
- 153 Zheng Zhang and Kil To Chong. Comparison between first-order hold with zero-order hold in
154 discretization of input-delay nonlinear systems. In *ICCAS*, pages 2892–2896, 2007.